CS/ME/ECE 759
High Performance Computing for Engineering Applications
Assignment 4
Due Thursday 10/6/2022 at 9:00 PM

Submit responses to all tasks which do not specify a file name to Canvas in a file called assignment4.{txt, docx, pdf, rtf, odt} (choose one of these formats). For this assignment, this means the two "beyond call of duty" problems that you may choose to do. Submit all plots on Canvas. Do not zip your Canvas submission.

All *source files* should be submitted in the `HW04` subdirectory on the `main` branch of your `GitLab` repo. Please use the name `HW04` exactly as shown here (both in terms of capitalization & name). The `HW04` subdirectory should have no subdirectories. For this assignment, your `HW04` folder should contain `task1.cu`, `task2.cu`, `matmul.cu`, and `stencil.cu`.

All commands or code must work on *Euler* with only the `nvidia/cuda` module loaded. Loading the module is done via
`$ module load nvidia/cuda`
Since various commands may behave differently on your computer, we recommend that you test on *Euler* before you submit your homework.

Please submit clean code. Consider using a formatter like clang-format.

IMPORTANT: Before you begin, copy any provided files from `Assignments/HW04` directory of the ME759 Resource Repo. Do not change any of the provided files since these files will be overwritten with clean, reference copies when grading.

---

1. (a) Implement in a file called `matmul.cu` the `matmul` and `matmul_kernel` functions as declared and described in the comment section of `matmul.cuh`. These functions should compute the product of square matrices.

   (b) Write a program `task1.cu` which will complete the following (some memory management steps are omitted for clarity, but you should implement them in your code for it to work properly):

   - Create matrices (as 1D row major arrays) `A` and `B` of size `n` $\times$ `n` on the host.
   - Fill these matrices with random numbers in the range [-1, 1].
   - Prepare arrays that are allocated as device memory (they will be passed to your `matmul` function.)
   - Call your `matmul` function.
   - Print the last element of the resulting matrix.
   - Print the time taken to execute your `matmul` function in *milliseconds* using CUDA events.
   - Compile: `nvcc task1.cu matmul.cu -Xcompiler -O3 -Xcompiler -Wall -Xptxas -O3 -std c++17 -o task1`
   - Run (where `n` and `threads_per_block` are positive integers): `./task1 n threads_per_block`
     - Note `n` is not necessarily a power of 2.
     - Use `Slurm` to run your job on Euler
   - Example expected output:
     `-16.35`
     `1.23`

   (c) On an Euler *compute node*, run `task1` for each value `n` $= 2^5, 2^6, \cdots, 2^{14}$ and generate a plot `task1.pdf` which plots the time taken by your algorithm as a function of `n` when `threads_per_block` $= 1024$. Overlay another plot which plots the same relationship with a different `threads_per_block` of your choice.

(d) Going beyond the call of duty, do if you wish to: Compare the scaling results with the results obtained in a previous assignment where you did a similar scaling analysis using a sequential implementation on the CPU. What do you see?

2. (a) Implement in a file called `stencil.cu` the `stencil` and `stencil_kernel` functions as declared and described in the comment section of `stencil.cuh`. These functions should produce the 1D convolution of `image` and `mask` as the following:

$$\texttt{output}[i] = \sum_{j=-R}^{R} \texttt{image}[i+j] * \texttt{mask}[j+R] \qquad i = 0, \cdots, \texttt{n} - 1 \ .$$

Assume that $\texttt{image}[i] = 1$ when $i < 0$ or $i > \texttt{n} - 1$. Pay close attention to what data you are asked to store and compute in shared memory.

(b) Write a program `task2.cu` which will complete the following (some memory management steps are omitted for clarity, but you should implement them in your code):

- Create arrays `image` (length `n`), `output` (length `n`), and `mask` (length `2 * R + 1`) on the host.
- Fill the `image` and `mask` array with random numbers in the range [-1, 1].
- Prepare arrays that are allocated as device memory (they will be passed to your `stencil` function.)
- Call your `stencil` function.
- Print the last element of the resulting `output` array.
- Print the time taken to execute your `stencil` function in *milliseconds* using CUDA events.
- Compile: `nvcc task2.cu stencil.cu -Xcompiler -O3 -Xcompiler -Wall -Xptxas -O3 -std c++17 -o task2`
- Run via `Slurm` (where `n`, `R`, and `threads_per_block` are positive integers): `./task2 n R threads_per_block`
- Example expected output:
  ```
  11.36
  1.23
  ```

(c) On an Euler *compute node*, run `task2` for each value $\texttt{n} = 2^{10}, 2^{11}, \cdots, 2^{29}$ and generate a plot `task2.pdf` which plots the time taken by your algorithm as a function of `n` when `threads_per_block` $= 1024$ and `R` $= 128$. Overlay another plot which plots the same relationship with a different `threads_per_block` of your choice.

(d) Going beyond the call of duty, do if you wish to: Compare the scaling results with the results obtained in a previous assignment where you did a similar scaling analysis using a sequential implementation on the CPU. What do you see?